

SyncTwin: Fast Digital Twin Construction and Synchronization for Safe Robotic Grasping

Anonymous CVPR submission

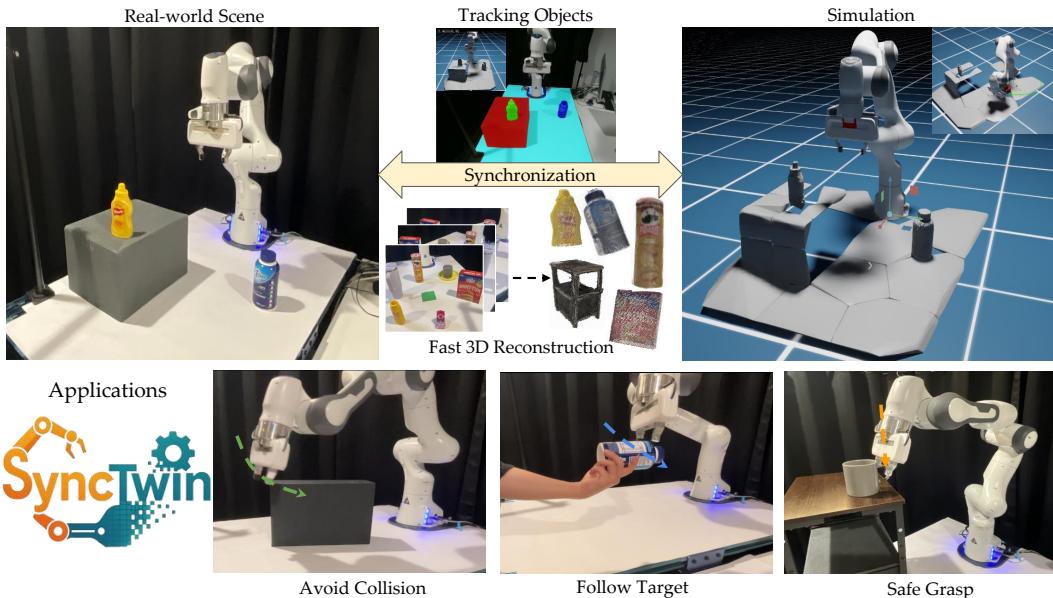


Figure 1. SyncTwin enables fast digital twin and synchronization based on efficient 3D scene reconstruction and object tracking, bridging motion planning for sim-to-real execution, which can be applied to dynamic obstacle avoidance, object tracking, and safe grasping under single-view occlusion in real-world environments.

Abstract

Accurate and safe grasping under dynamic and visually occluded conditions remains a core challenge in real-world robotic manipulation. We introduce SyncTwin, a novel digital twin framework that unifies fast 3D scene reconstruction and real-to-sim synchronization for robust and safety-aware grasping in such environments. In the offline stage, we employ VGGT to rapidly reconstruct object-level 3D assets from RGB images, forming a reusable geometry library for simulation. During execution, SyncTwin continuously synchronizes the digital twin by tracking real-world object states via point cloud segmentation updates and aligning them through colored-ICP registration. The updated twin enables motion planners to compute collision-free and

dynamically feasible trajectories in simulation, which are safely executed on the real robot through a closed real-to-sim-to-real loop. Experiments in dynamic and occluded scenes show that SyncTwin improves grasp accuracy and motion safety, demonstrating the effectiveness of digital-twin synchronization for real-world robotic execution.

1. Introduction

Achieving accurate and safe robotic grasping in dynamic real-world environments remains a long-standing challenge, due to incomplete perception and dynamic scenes. Without an accurate understanding of their surroundings, robots risk colliding with the environment, which may damage hard-

014
015
016
017
018
019

020
021
022
023
024
025

026 ware or even endanger humans [58]. Thus, ensuring that
027 robots can safely plan and execute motions amid dynamic
028 scene changes is a prerequisite for reliable real-world de-
029 ployment. Unlike simulation, where complete scene geom-
030 etry and object states are fully accessible, real-world per-
031 ception often only offers partial, occluded observations. As
032 a result, motion-planning algorithms [16, 33, 36] that rely
033 on full and accurate environment models in simulation face
034 significant challenges when deployed in the real world.

035 Several recent efforts have attempted to bridge the gap
036 between perception and control. Voxel-based mapping sys-
037 tems like NVBlox [23] enable online obstacle reconstruc-
038 tion, yet they only provide a voxel grid, which is a represen-
039 tation too coarse for reliable manipulation. End-to-end re-
040 active policies like DRP [53] handle dynamics through con-
041 tinuous control but require robot-specific retraining and lack
042 generalization across new hardware or environments. Thus,
043 all of these methods share a critical limitation: they oper-
044 ate without a consistent and complete model of the scene,
045 leading to unsafe or unreliable executions in the real world.

046 Bridging this gap between simulation and the real world,
047 a robot needs a perception model that can efficiently per-
048 ceive what objects exist in the environment and also track
049 how the real-world scene evolves over time. In other words,
050 instead of planning in a static or outdated simulation, the
051 robot should plan within a dynamic digital twin that mirrors
052 the physical world in real time, where perception and con-
053 trol are tightly coupled through continuous real-to-sim syn-
054 chronization. However, real-world perception is inherently
055 partial—single-view occlusions often reveal only fragments
056 of object geometry, making grasping and motion planning
057 unreliable. Inspired by SAM4D [51], which maintains a
058 memory bank of object assets, we incorporate the idea of
059 leveraging object-level memories to complete partial obser-
060 vations at execution time. We develop **SyncTwin**, a digital-
061 twin framework equipped with an object memory bank that
062 performs real-time object tracking from point clouds, in-
063 jects accurate poses and geometries into simulation, and
064 closes the loop by executing planned trajectories back on
065 the real robot, as illustrated in Figure 1.

066 SyncTwin operates in two stages. In Stage I, VGGT [46]
067 generates scene-level point clouds from RGB inputs.
068 Object-level point clouds are extracted via projection, seg-
069 mentation, and denoising, then converted into lightweight
070 meshes and stored in a memory bank. In Stage II, the digital
071 twin is continuously synchronizes with the real world using
072 SAM2 for object tracking and GPU-accelerated colored-
073 ICP [42]. Combined, these two methods align partial obser-
074 vations with stored assets to maintain a consistent and
075 complete scene representation. The updated digital twin
076 is streamed into Isaac Sim [24], allowing cuRobo [41] to
077 perform motion planning and generate collision-aware tra-
078 jectories. Moreover, this architecture can be transferred to

079 different real robots without any retraining.

080 The main contributions of this paper are as follows:

- We present the first digital twin framework that tracks 3D objects in real time from point clouds and updates their poses and geometries in a synchronized simulation, enabling collision-aware planning and a closed real-to-sim-to-real loop for dynamic, partially observed scenes.
- We introduce a fast, low-cost RGB-only method that constructs 3D geometry assets using learning-based geometry estimation and projection-based segmentation.
- We develop a real-time 3D segmentation and tracking module that processes streaming RGB-D camera data.
- Our experiments demonstrate that the proposed system improves grasp accuracy and safety in single-view and occluded settings, while also achieving state-of-the-art efficiency in 3D geometry asset reconstruction.

2. Related Work

095 **3D Scene Reconstruction and Segmentation.** Traditional
096 3D scene reconstruction approaches rely on RGB-D in-
097 put from depth sensors to estimate geometry [5, 7, 28,
098 38]. More recent approaches operate purely on RGB im-
099 ages, either through optimization-based pipelines such as
100 Structure-from-Motion (SfM) [14, 30, 31, 35] and Multi-
101 View Stereo (MVS) [10, 11, 37], or through learning-based
102 frameworks such as MVSNet [54], NeRF [22], and 3D
103 Gaussian Splatting [20]. While these methods can recover
104 visually plausible geometry, they often require substantial
105 computation time to produce complete reconstructions. For
106 robotics applications, object-centric 3D segmentation is im-
107 portant as well. Prior work has explored segmentation di-
108 rectly on point clouds using learned clustering and aggre-
109 gation strategies [15, 47, 57], as well as multi-view projec-
110 tion of 2D masks for 3D instance segmentation [21, 34, 56].
111 However, projection-based methods typically require accu-
112 rate camera extrinsics [2, 52], and little research has ex-
113 amined how such strategies perform when applied to point
114 clouds produced by learning-based reconstruction methods.
115 Our work addresses this gap by unifying learned 3D recon-
116 struction with mask-guided multi-view segmentation, en-
117 abling efficient and robust object asset generation for down-
118 stream robotic manipulation.

119 **Digital Twin for Robotic Manipulation.** Digital twin
120 systems have become increasingly popular for sim-to-real
121 transfer, particularly in training reinforcement learning and
122 imitation learning policies [4, 45, 50, 55]. And several
123 studies leverage generated digital twins as a form of data
124 augmentation to enhance the generalization of downstream
125 models [19, 25]. However, existing frameworks typically
126 reconstruct static scenes once before training [29] and do
127 not maintain continuous synchronization with the evolving
128 physical world. Some recent efforts attempt real-time track-
129 ing by detecting object locations with 2D detectors [40], but

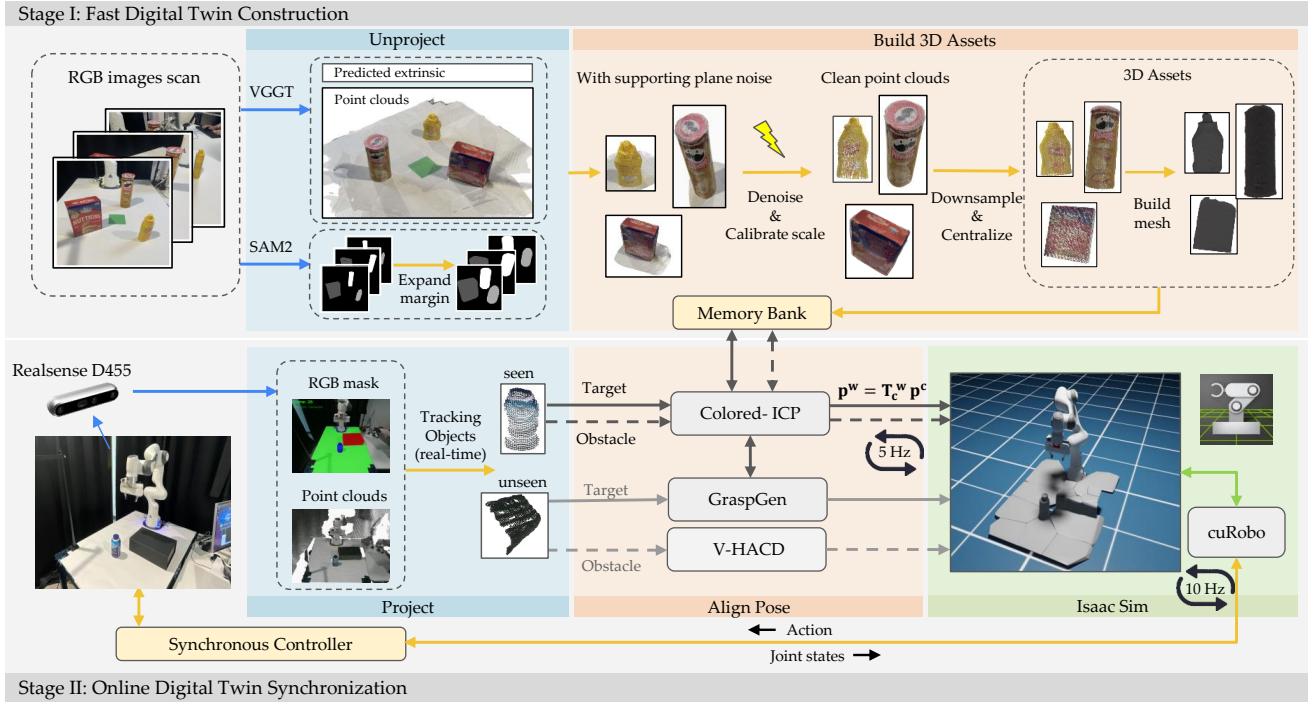


Figure 2. Framework of the SyncTwin. Stage I reconstructs simulation-ready 3D assets from RGB images using VGGT and SAM2. Multi-view masks are unprojected into point clouds, then denoised, scaled, and meshed into clean object assets stored in a memory bank. Stage II performs real-time object segmentation, pose tracking, and asset-based completion, enabling grasp generation and reactive motion planning in a closed real2sim2real loop. By continuously updating the digital twin and leveraging simulation for decision making, the system ensures safe and adaptive execution under dynamic and partially occluded environments.

they often lack precise pose estimation or geometry updates. In contrast, our system enables continuous synchronization of the digital twin with real-world online perception, which enables accurate tracking and reliable manipulation in dynamic environments under occlusion.

Safe Motion Planning for Manipulation. Safe and robust motion planning remains a critical challenge for robotic manipulation in unstructured environments [18, 39]. Classical and learning-based planners [16, 33, 36] typically operate in simulation, where the environment is fully accessible. When deployed in the real world, however, they must handle partial observations. To narrow this gap, several methods construct static point cloud maps and import them into simulation for offline planning [3], though such maps struggle to support real-time adaptation. Other approaches aim to ensure safety by predicting collision-free actions directly from images in latent space [1, 27], using force-sensing-based control [17, 48], or distilling planning policies from point cloud observations [6, 9, 53]. NVBlox improves online safety by voxelizing scene geometry and segmenting the robot in real time [23]. In contrast, our method maintains a dynamically synchronized digital twin that continuously provides updated scene geometry to the planner, enabling safe execution in dynamic, cluttered environments.

3. Method

SyncTwin consists of two stages: (1) fast digital twin construction, and (2) digital twin synchronization. An overall framework is provided in Figure 2.

3.1. Problem Formulation

We aim to enable safe robotic grasping in dynamic, partially observable real-world environments by maintaining a *continuously synchronized digital twin*. This problem can be decomposed into the following components:

Stage I: Fast Digital Twin Construction. The system receives a small set of RGB images $\{\mathbf{I}_i\}_{i=1}^N$ along with camera intrinsics K and estimated extrinsics $\{\mathbf{T}_i^{\text{world}}\}_{i=1}^N$. The goal is to produce object-level, simulation-ready 3D assets $\mathbb{B} = \{\mathcal{X}_j, \mathcal{M}_j^{3D}\}$ from these images. The main challenge is that learning-based extrinsics contain unstable errors, causing mask–projection misalignment and table–object mixing in the reconstructed point cloud, which must be addressed to obtain clean object geometry.

Stage II: Online Digital Twin Synchronization. During execution, the system receives streaming RGB-D frames and corresponding partial point clouds \mathcal{X}_p . The objective is to maintain accurate object poses $\mathbf{T}_j^{\text{world}}$ in the simula-

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177 tor by aligning \mathcal{X}_p with their complete assets $\mathcal{X}_m \in \mathbb{B}$.
 178 This synchronized scene is streamed into Isaac Sim [24],
 179 where cuRobo’s MPC planner [41] produces short-horizon,
 180 collision-free trajectories $\mathbf{A}_{t:t+H} = \{\mathbf{a}_0, \dots, \mathbf{a}_H\}$. The
 181 key challenge is robustly tracking objects under occlusion
 182 and partial observation, while ensuring that the object poses
 183 and geometries can be accurately updated into the simulator
 184 to enable safe real-to-sim-to-real planning.

185 3.2. Fast Digital Twin Construction

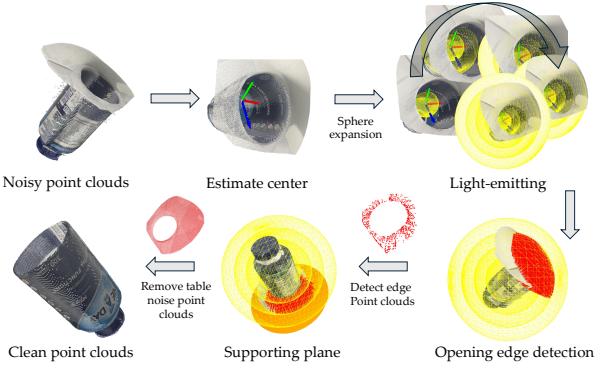
186 The first stage aims to rapidly reconstruct the 3D environment
 187 and extract object-level representations suitable for
 188 real-time simulation, where the focus is to achieve accurate
 189 geometric perception for motion planning, rather than photo-
 190 realistic reconstruction. We employ VGGT [46] to recon-
 191 struct dense scene point clouds directly from a small num-
 192 ber of RGB images, which enables efficient extraction of
 193 object-level geometry in a fast and low-cost manner with-
 194 out depth sensors or multi-view optimization.

195 Nevertheless, there are two major practical challenges.
 196 First, the camera extrinsics estimated by VGGT are often
 197 inaccurate, which leads to projection misalignment dur-
 198 ing mask-based segmentation. Second, the generated point
 199 cloud is not in a true metric world scale, causing the im-
 200 ported assets to appear incorrectly sized relative to the robot
 201 in the simulator. Therefore, Stage I focuses on producing
 202 simulation-ready 3D digital assets from VGGT outputs by
 203 addressing these limitations. We design a four-step pipeline
 204 that corrects extrinsic inaccuracies, enforces scale consis-
 205 tency, and generates clean object meshes suitable for down-
 206 stream planning, which is introduced as follows.

207 **Mask Projection Expansion.** To mitigate inaccuracies in
 208 VGGT-estimated camera extrinsics, each 2D segmentation
 209 mask \mathcal{S}_i is spatially expanded before projection, which
 210 ensures full coverage of object boundaries and prevents miss-
 211 ing edge regions during 3D reconstruction. While mask ex-
 212 pansion compensates for projection drift, it also introduces
 213 floating outliers (e.g., background points above the object)
 214 and merged support-plane regions (e.g., table surfaces). To
 215 address this, we apply our point clouds denoising mecha-
 216 nism to isolate the true object shape.

217 **Point Clouds Denoising.** Detecting openings or cavities in
 218 3D point clouds is a fundamental step in shape understand-
 219 ing for denoising table noise. We propose a purely geom-
 220 etric method that progressively expands a virtual light sphere
 221 from the object’s center and tracks uncovered regions on the
 222 spherical sampling space. The algorithm automatically de-
 223 tects openings and extracts rim points around their bound-
 224 aries without requiring mesh topology or prior segmenta-
 225 tion. Figure 3 shows the overall process.

226 Given a point cloud \mathcal{P} , we first estimate a geometric
 227 center $\mathbf{c} = \text{mean}(\mathcal{P})$. We then discretize the unit sphere
 228 into F directions $\{\mathbf{d}_i\}_{i=1}^F$ using a Fibonacci spiral distri-



253 **Figure 3. Supporting-plane noise removal mechanism.** A vir-
 254 tual light sphere expands from the object center to identify open-
 255 ings and boundary points, enabling filtering of table-plane noise.

256 bution [13], forming a uniform sampling domain \mathcal{D} . Each
 257 point \mathbf{p}_i defines a normalized direction $\hat{\mathbf{v}}_i = (\mathbf{p}_i - \mathbf{c}) / \|\mathbf{p}_i - \mathbf{c}\|$ and is assigned to its nearest angular bucket direction \mathbf{d}_j if $\hat{\mathbf{v}}_i \cdot \mathbf{d}_j \geq \cos(\theta_{\text{tolerance}})$. We then iteratively expand a vir-
 258 tual sphere centered at \mathbf{c} with radius r_t :

$$r_{t+1} = r_t + \Delta r, \quad r_{\min}(j) = \min_{i \in \text{bucket } j} \|\mathbf{p}_i - \mathbf{c}\|. \quad (1)$$

259 A direction \mathbf{d}_j is marked as *hit* once any assigned point en-
 260 ters the sphere, and unhit directions form a binary mask \mathcal{U}_t .
 261 Stable uncovered regions (openings or cavities) are detected
 262 when the largest unhit component remains consistent over
 263 iterations. We denote by $\mathcal{N}(j) = \{k \mid (j, k) \in \mathcal{P}\}$ the
 264 neighborhood of bucket j . Boundary buckets are defined as
 265 unhit directions adjacent to hit ones:

$$\mathcal{B} = \{j \in \mathcal{U} \mid \exists k \in \mathcal{N}(j), \text{hit}(k) = 1\}. \quad (2)$$

266 For each boundary \mathbf{d}_j , the farthest point within tolerance
 267 is chosen as a rim sample \mathbf{p}_j^* . Connected components
 268 on the spherical adjacency graph are extracted to identify
 269 large uncovered regions. Let $\{\mathcal{C}_m\}$ denote all connected
 270 components of the unhit set \mathcal{U}_t , \mathcal{C}_{\max} is the largest con-
 271 nected region. The principal opening direction is then com-
 272 puted by averaging the largest unhit component: $\mathbf{n}_{\text{open}} =$

273
$$\frac{\sum_{j \in \mathcal{C}_{\max}} \mathbf{d}_j}{\|\sum_{j \in \mathcal{C}_{\max}} \mathbf{d}_j\|}$$
. Finally, the rim points $\{\mathbf{p}_j^*\}$ are fitted with
 274 a plane using SVD [12], yielding the opening orientation
 275 and visualizable boundary. Compared with RANSAC plane
 276 fitting [8], which can only segment dominant planes, our
 277 method can detect cavity openings, enabling the identifica-
 278 tion of ring-shaped planes surrounding the opening

279 **Real-World Scale Alignment.** Since VGGT produces
 280 point clouds that are not in a world-scale metric, we esti-
 281 mate a global scale factor to align the reconstructions with
 282 world coordinates. Even with known intrinsic parameters,
 283 according to the pinhole camera model, monocular geom-
 284 etry cannot determine absolute scale [14]. Therefore, we

calibrate the scale using a reference object or markers of known physical dimensions within the scene. Implementation details are provided in supplementary.

Mesh Simplification. To maintain real-time performance in the digital twin, we apply an adaptive mesh decimation that reduces vertex count while preserving geometric fidelity and collision boundaries. To avoid over-smoothing across sharp edges, we use an angle-based gating weight $w_{ij} = 1$ if $\theta_{ij} \leq \theta_{\text{th}}$ and 0 otherwise, where θ_{th} is a feature threshold (e.g., 30°) and $\theta_{ij} = \arccos(\mathbf{n}_i^\top \mathbf{n}_j)$ for all $j \in \mathcal{N}(i)$. Each vertex is then updated via a selective Laplacian step [44]. Compared to uniform mesh decimation, this feature-aware smoothing preserves sharp edges around handles and object rims, which are critical for accurate grasp planning. All processed point clouds, meshes, and their id are stored in a memory bank, which serves for object recognition and scene synchronization in Stage II.

3.3. Online Digital Twin Synchronization

The second stage of our system focuses on real-time object tracking and safe grasp execution through continuous perception, planning synchronization between the real and digital environments. This stage consists of four tightly integrated modules: real-time point cloud segmentation, GPU-accelerated colored-ICP registration, grasp pose generation from complete object models, and dynamic motion planning with cuRobo MPC.

Real-time Point Clouds Segmentation. To achieve real-time segmentation on incoming RGB-D streams, we build a module shown in Fig 4, that performs continuous inference on camera frames and outputs segmentation masks \mathcal{S}_p , and then projects the mask onto the full point cloud to obtain the corresponding partial object point clouds \mathcal{X}_p . Compared with traditional offline SAM2 inference, we design a sliding window mechanism that enables SAM2 to process camera streams in real time, maintaining temporal consistency in object masks for continuous 3D segmentation and tracking under occlusions. At image size 640×480 , our method runs at 15 Hz on RTX 4090.

Colored-ICP Registration. For aligning partial point clouds \mathcal{X}_p obtained from the camera with the corresponding full object model \mathcal{X}_m stored in the memory bank, we employ a colored-ICP algorithm [32] implemented on the GPU via the cupoch library [42]. Unlike traditional geometric ICP, which minimizes only spatial distance, colored-ICP jointly minimizes geometric and color residuals: $E(R, t) = \sum_i [\lambda_g \|R\mathbf{x}_i + t - \mathbf{y}_i\|^2 + \lambda_c \|I(\mathbf{x}_i) - I(\mathbf{y}_i)\|^2]$ where (R, t) denotes the transformation between \mathcal{X}_p and \mathcal{X}_m , and $I(\cdot)$ represents point color intensity. The weighting factors λ_g and λ_c balance geometric and color terms.

Grasp Pose Generation and Obstacle Representation. For unseen target objects, we directly apply GraspGen to the partial point cloud. In contrast, once a seen target ob-

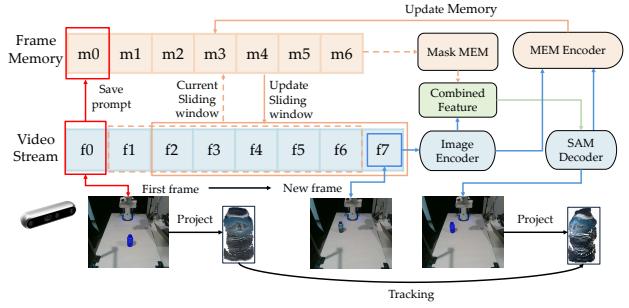


Figure 4. **Overview of the camera predictor module.** The red solid line indicates that the first frame is persistently stored in the frame memory. The yellow dashed line represents the operation of saving frames from the previous time step into the memory. The blue and green solid lines denote the data flow and processing steps for the current frame. Together, the sliding-window mechanism enables real-time video segmentation and object-level point cloud tracking with temporally consistent memory updates.

ject is registered, we replace its partial observation with the complete point cloud \mathcal{X}_m from the memory bank and feed it into GraspGen to predict grasp poses $\{\mathbf{T}_{\text{gripper}} = f_{\text{GraspGen}}(\mathcal{X}_m)\}$. This replacement mitigates the uncertainty from occlusion and single-view perception, yielding more stable and accurate grasp pose estimation. For unseen obstacles, we dynamically generate multi-convex hulls (V-HACD [49]) from the segmented point cloud for collision modeling. For known obstacles, the aligned object meshes and poses are directly imported into the digital twin for real-time collision checking.

Motion Planning with Sim-to-real Synchronization. Finally, motion planning and control are performed using the cuRobo’s model predictive control (MPC) framework. At each control step, the robot’s joint states are synchronized with the digital twin, where cuRobo computes an optimized short-horizon trajectory under real-time collision constraints. Only the first control action $\{\mathbf{a}_0\}$ from the predicted trajectory $\mathbf{A}_{t:t+H} = \{\mathbf{a}_0, \mathbf{a}_1, \dots, \mathbf{a}_H\}$ is executed on the real robot, followed by continuous replanning with the environment updates, achieving closed-loop synchronization between simulation and reality.

4. Experiments

We evaluate SyncTwin across three dimensions: (1) the efficiency of the fast 3D reconstruction pipeline; (2) obstacle avoidance performance under dynamic and single-view occluded conditions; and (3) grasp success rate under single-view occlusion. All experiments are designed to validate both the offline and online components of our framework.

4.1. Experiment Setup

All experiments are conducted with a Franka Emika Panda robotic arm equipped with an Intel RealSense D455 RGB-

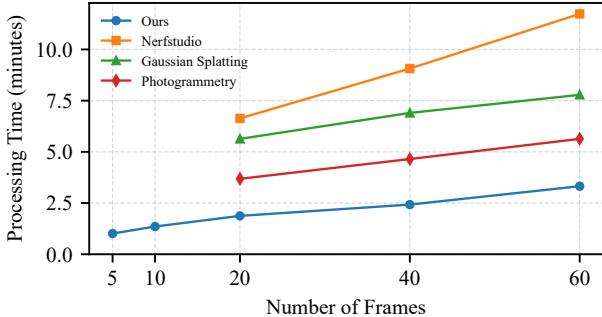


Figure 5. **The comparison of processing time given different numbers of input images.** The processing time covers both reconstruction and segmentation. The 5 and 10 images do not apply to the baselines because of fail to estimate the camera extrinsic.

D camera mounted above and in front of the workspace. And iPhone 12 for RGB images. We apply voxel-based downsampling to the input point cloud with a voxel size of 3 mm. This is motivated by the fact that the computational complexity of point cloud segmentation algorithms scales linearly with the number of points. The digital twin is implemented in Isaac Sim 4.0 and integrated with the cuRobo MPC motion planning framework, running on a single NVIDIA RTX 4090 GPU. Both perception and planning run on the same GPU, enabling a closed-loop update rate of up to 5 Hz. Motion planning runs at 10 Hz, and the robot’s velocity scaling is set to 0.2. The test objects include bottles, cans, cups, and boxes of various shapes.

4.2. Baselines and Metrics

Baselines. (1) For *3D reconstruction*: we compare against Photogrammetry [35], NeRF (Nerfstudio [43]), and Gaussian Splatting (3DGs [20]). (2) For *obstacle avoidance*: we adopt NVBlox [23] as the baseline voxel-mapping method. For the ablation studies, we analyze performance across the following settings: (1) *Mask Expansion and Denoising*: using variants without mask expansion and without denoising as baselines, evaluating their reconstruction quality against our full segmentation mechanism. (2) *Object Completion*: we use GraspGen [26] as the baseline grasp generator, comparing grasp poses predicted from single-view partial point clouds (baseline) versus from complete, asset-retrieved geometry produced by SyncTwin (ours).

Evaluation Metrics. Depending on the experiment, we report the following metrics: reconstruction time (min), dependency on the number of input images N_{\min} , obstacle avoidance success rate (%), and grasp success rate (%). For avoidance tests, we define 3-level outcome levels: FA (full avoidance, no contact with the object, 1.0), EA (edge avoidance, slight contact without displacing the object, 0.8), and CO (collision, noticeable contact that significantly moves

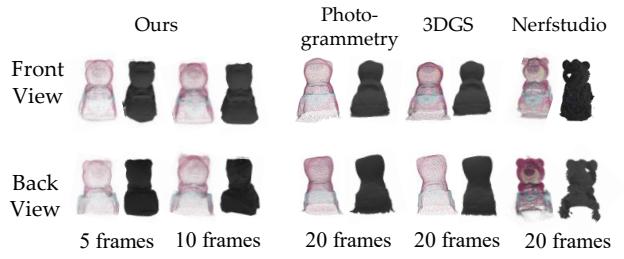


Figure 6. **Reconstruction comparison given different numbers of input images.** In each column, the left figure shows the point clouds, and the right one shows the untextured mesh.

the object, 0.0). Weighted success rate:

$$SR = \frac{N_{FA} + 0.8 N_{EA} + 0.0 N_{CO}}{N} \times 100\%.$$

4.3. Fast 3D Reconstruction

We evaluate the efficiency and input-image dependency of SyncTwin’s 3D reconstruction module, comparing against Photogrammetry, 3DGs, and Nerfstudio. Each method reconstructs the same object using 5, 10, 20, 40, 60 RGB images, and all approaches obtain object assets through multi-view projection based segmentation. The experimental results are summarized as follows.

Reconstruction Time. As shown in Figure 5, SyncTwin achieves the shortest reconstruction time across all settings. It produces a simulation-ready mesh in only about 1–2 minutes using 5–10 input images, while Photogrammetry, 3DGs, and Nerfstudio require at least 4–7 minutes even with more frames. This enables significantly faster digital-twin construction in a new real-world scenarios.

Dependency on the Number of Input Images. SyncTwin also exhibits the lowest dependency on input-image count. It generates usable meshes from as few as 5–10 images, whereas competing approaches typically require 20+ images to avoid failure caused by unstable optimization of camera extrinsics. This low image dependency substantially accelerates asset generation and improves robustness under limited views, especially because the reported processing time excludes image-capture time, which becomes additional and unpredictable when more images are required.

Furthermore, the qualitative comparison in Figure 6 shows that our method preserves fine-grained geometric details even with very few input images, as our mesh simplification algorithm maintains high geometric fidelity while reducing mesh vertices, enabling faster simulation performance. For instance, the *bear’s ear shape* remains well preserved with only 5–10 frames, demonstrating both low image dependency and strong geometric consistency.

These properties make SyncTwin a state-of-the-art solution for fast 3D geometry asset generation: the system

Table 1. Comparison of obstacle-avoidance performance between NVBlox and SyncTwin in dynamic environments. Unseen objects (left) are not present in the asset memory, whereas starred objects in the Seen category (right) are stored in the memory bank. SyncTwin achieves significantly higher success rates in both settings, with particularly strong gains when complete object assets are available.

Method	Motion	Unseen												Seen (Memory Bank)													
		Box1			Box2			Box3			Box4			SR_unseen (%)	Box1*			Box2*			Box3*			Box4*			SR_seen (%)
		FA	EA	CO	FA	EA	CO	FA	EA	CO	FA	EA	CO		FA	EA	CO										
NVBlox	SelfRot	6	4	10	9	5	6	3	4	13	3	11	6	50.3%	—	—	—	—	—	—	—	—	—	—	—	—	—
	EnterTraj	3	9	8	4	9	7	1	6	13	0	3	17	37.0%	—	—	—	—	—	—	—	—	—	—	—	—	—
SyncTwin	SelfRot	12	7	1	14	5	1	11	7	2	9	9	2	85.5%	18	1	1	16	4	0	15	5	0	13	6	1	93.5%
	EnterTraj	8	8	4	11	6	3	8	7	5	7	8	5	71.5%	12	7	3	12	5	3	10	6	4	9	7	4	78.8%

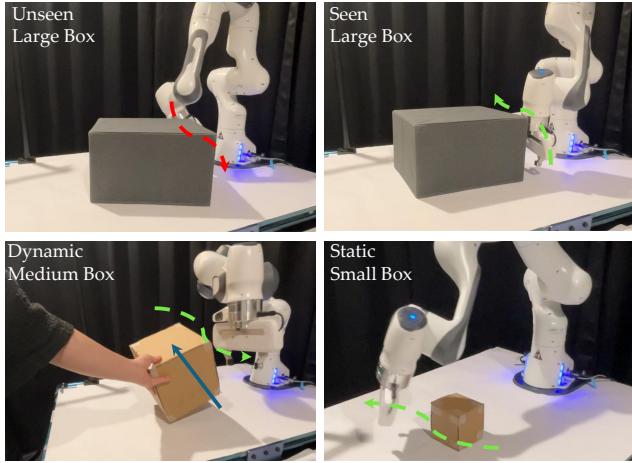


Figure 7. **Examples of SyncTwin’s dynamic obstacle avoidance.** Green dashed lines indicate collision-free robot trajectories, while red lines mark trajectories that result in collisions. Blue arrows denote the motion of dynamic obstacles. For unseen objects (top-left), the robot collides with unobserved regions, whereas the same object in the seen case (top-right) is successfully avoided. SyncTwin also handles dynamic obstacles (bottom-left) and small objects (bottom-right) effectively.

can construct high-quality, simulation-ready object meshes in about 1 minute, enabling rapid reconstruction of object assets for downstream digital-twin synchronization.

4.4. Obstacle Avoidance under Occlusion

We adopt the built-in obstacle avoidance benchmark in cuRobo to evaluate the success rate of obstacle avoidance, where the robot repeatedly moves between two target points while avoiding obstacles along its path.

Both NVBlox and SyncTwin are tested on unseen (not stored in memory bank) obstacles to evaluate avoidance performance. Additionally, SyncTwin is evaluated on seen (stored in memory bank) objects recorded in its asset memory to study how prior geometry improves safety and reactivity under single-view occlusion. The experiment tests under two motion patterns: **SelfRot** indicates in-place rota-

tion; **EnterTraj** blocks motion into the predicted trajectory. Each condition is repeated for 20 trials, total $N = 20 \times 4 = 80$. We test four boxes with sizes of Box1: $10 \times 10 \times 10$ cm, Box2: $20 \times 20 \times 20$ cm, Box3: $10 \times 20 \times 30$ cm, and Box4: $20 \times 22 \times 35$ cm as obstacles. Through experimental observation, the results are as follows:

Unseen Object Performance. As shown in Table 1, SyncTwin consistently outperforms NVBlox when encountering unseen obstacles under single-view dynamic occlusion. NVBlox exhibits failure modes such as: (1) unstable voxelization for small objects (e.g., Box1) due to sparse depth returns; (2) misclassification of limited height objects (e.g., Box3) as part of the tabletop caused by inconsistent depth estimation; and (3) trajectory intersections when large obstacles (e.g., Box4) fall outside the sensor’s visible region. Across all these cases, SyncTwin maintains markedly higher obstacle-avoidance success rates, demonstrating stronger robustness to object scale, occlusion, and limited viewpoint coverage.

Seen Object Performance. In Table 1, when objects are stored in the asset memory, SyncTwin’s performance improves even further. Complete geometric priors obtained in Stage I resolve the challenges posed by small or thin objects, eliminate many failure modes, and convert numerous edge-avoidance cases into full avoidance. These results indicate that SyncTwin not only generalizes better on unseen objects but also achieves substantially higher reliability when the objects have been previously reconstructed. These results highlight the importance of SyncTwin’s Stage I memory construction, which enables complete geometric reasoning even under partial observability. Furthermore, Figure 7 illustrates representative avoidance examples achieved by SyncTwin.

4.5. Ablation Studies

Mask Expansion and Denoising. Due to the inaccuracy of VGGT-estimated camera extrinsics, multi-view mask projections often become misaligned and fail to fully cover the object. Our ablation study evaluates reconstruction results *w/o mask expansion* and *w/o denoising*. As shown in

433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471

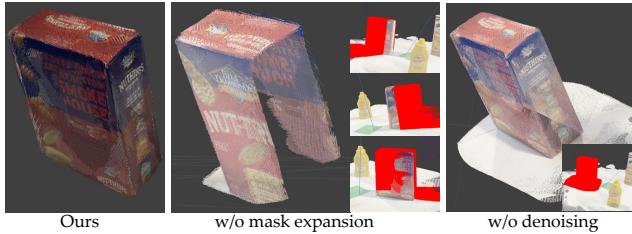


Figure 8. **Comparison of segmentation-based 3D reconstruction results.** Left: Our method generates a complete and clean object mesh without support-plane noise. Middle: w/o mask expansion, multi-view projections (red) fail to fully overlap, and their intersection leads to missing object geometry. Right: w/o denoising, the intersection of projected regions (red) incorrectly preserves support-plane points, introducing significant noise into the reconstructed point cloud.

Table 2. **Comparison of grasp success rates before and after geometry completion.** Asset-based completion significantly improves performance across all objects by producing more accurate and safer grasp pose candidates, with the largest gains observed for partially occluded items such as the cup (with handle).

Condition	Bottle	Cup (Handle)	Cookie Box	Chips Can
Before Completion (%)	78.3	65.0	81.7	80.0
After Completion (%)	90.0	86.7	93.3	95.0
Improvement	+11.7	+21.7	+11.6	+15.0

Figure 8, removing margin expansion causes inconsistent projections across viewpoints, leaving parts of the object missing. Likewise, disabling supporting-plane points denoising preserves table artifacts in the point cloud, resulting in noisy meshes. These comparisons demonstrate that segmentation-aware expansion and support-plane denoising are essential for obtaining clean and stable 3D assets.

Object Completion. We evaluate the effect of object completeness on grasp generation. When grasp poses are generated from single-view partial point clouds, the limited geometry often leads to incomplete or incorrect grasp configurations, for example. As shown in Figure 9, the cup handle is misinterpreted, resulting in unsafe or colliding grasps. In contrast, after SyncTwin aligns the observed object with the corresponding complete mesh from the asset library, the grasp generator produces denser, more accurate, and physically feasible grasp candidates. This demonstrates that asset-based completion is essential for reliable and collision-free grasp execution in the real world. As shown in Table 2, across 60 evaluation trials, grasp success rates increase substantially after geometry completion. The improvement is most pronounced for partially occluded or asymmetric objects, such as the cup and chips can, where asset-based completion provides the missing structure needed for reliable planning.

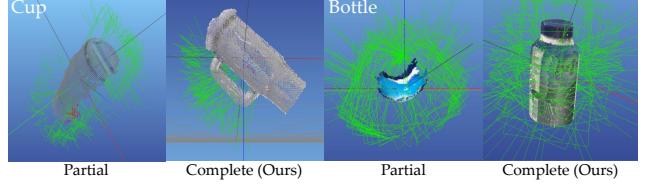


Figure 9. **Comparison of grasp candidates generated from Ours (w/o Completion) and Ours.** The example objects include a handled cup (left pairs) and a bottle (right pairs). Green lines visualize predicted gripper poses.

4.6. Discussions

In our experiments, Stage I reconstruction occasionally fails when the object lacks a clear supporting surface or when the support plane is weakly connected to the object (e.g., a tall cup). In these cases, inaccurate camera extrinsics cause the 2D segmentation masks to be incorrectly projected into 3D, resulting in broken or incomplete object point clouds and missing geometry. Such extrinsic errors produced by learning-based reconstruction remain an open challenge.

5. Conclusion

We introduced SyncTwin, a digital-twin framework that unifies fast RGB-only 3D reconstruction with real-time scene synchronization for safe and robust grasping in dynamic, partially occluded environments. By leveraging VGGT-based reconstruction, segmentation-aware denoising, and memory-driven geometry completion, SyncTwin provides accurate object geometry and reliable grasp generation from limited visual input. By bridging the sim-to-real gap through consistent digital-twin updates, the system enables simulation-based planners to execute safe, collision-aware trajectories on real robots without retraining.

Our experiments demonstrate substantial improvements over existing baselines such as NVBlox, including higher obstacle-avoidance success rates, more stable behavior in a dynamic environment, and improved grasp performance under single-view occlusion. These results highlight the advantage of combining fast asset generation with persistent memory and real-time synchronization, enabling safe and accurate execution in real-world environments.

Looking forward, two extensions appear particularly promising. First, enabling online asset expansion by integrating Stage I construction into Stage II synchronization would allow the system to incrementally acquire new objects as they appear. Second, distributed or multi-GPU system designs that decouple perception, synchronization, and planning could further improve responsiveness under fast dynamic scenes. Consequently, these enable more adaptive, real-world-oriented digital-twin synchronization.

535 References

- [1] Arpit Bahety, Arnav Balaji, Ben AbbateMatteo, and Roberto Martín-Martín. Safemimic: Towards safe and autonomous human-to-robot imitation for mobile manipulation. *arXiv preprint arXiv:2506.15847*, 2025. 3
- [2] Mohamed El Amine Boudjoghra, Angela Dai, Jean Lahoud, Hisham Cholakkal, Rao Muhammad Anwer, Salman Khan, and Fahad Shahbaz Khan. Open-yolo 3d: Towards fast and accurate open-vocabulary 3d instance segmentation. *arXiv preprint arXiv:2406.02548*, 2024. 2
- [3] Lukas Brunke, Yanni Zhang, Ralf Römer, Jack Naimer, Nikola Staykov, Siqi Zhou, and Angela P Schoellig. Semantically safe robot manipulation: From semantic scene understanding to motion safeguards. *IEEE Robotics and Automation Letters*, 2025. 3
- [4] Yevgen Chebotar, Ankur Handa, Viktor Makoviychuk, Miles Macklin, Jan Issac, Nathan Ratliff, and Dieter Fox. Closing the sim-to-real loop: Adapting simulation randomization with real world experience. In *2019 International Conference on Robotics and Automation (ICRA)*, pages 8973–8979. IEEE, 2019. 2
- [5] Angela Dai, Matthias Nießner, Michael Zollhöfer, Shahram Izadi, and Christian Theobalt. Bundlefusion: Real-time globally consistent 3d reconstruction using on-the-fly surface reintegration. *ACM Transactions on Graphics (ToG)*, 36(4):1, 2017. 2
- [6] Murtaza Dalal, Jiahui Yang, Russell Mendonca, Youssef Khaky, Ruslan Salakhutdinov, and Deepak Pathak. Neural mp: A generalist neural motion planner. *arXiv preprint arXiv:2409.05864*, 2024. 3
- [7] Mingsong Dou, Li Guan, Jan-Michael Frahm, and Henry Fuchs. Exploring high-level plane primitives for indoor 3d reconstruction with a hand-held rgb-d camera. In *Asian Conference on Computer Vision*, pages 94–108. Springer, 2012. 2
- [8] Martin A Fischler and Robert C Bolles. Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography. *Communications of the ACM*, 24(6):381–395, 1981. 4
- [9] Adam Fishman, Adithyavairavan Murali, Clemens Eppner, Bryan Peele, Byron Boots, and Dieter Fox. Motion policy networks. In *conference on Robot Learning*, pages 967–977. PMLR, 2023. 3
- [10] Yasutaka Furukawa and Jean Ponce. Accurate, dense, and robust multiview stereopsis. *IEEE transactions on pattern analysis and machine intelligence*, 32(8):1362–1376, 2009. 2
- [11] Silvano Galliani, Katrin Lasinger, and Konrad Schindler. Massively parallel multiview stereopsis by surface normal diffusion. In *Proceedings of the IEEE international conference on computer vision*, pages 873–881, 2015. 2
- [12] Gene H Golub and Christian Reinsch. Singular value decomposition and least squares solutions. In *Linear algebra*, pages 134–151. Springer, 1971. 4
- [13] Álvaro González. Measurement of areas on a sphere using fibonacci and latitude-longitude lattices. *Mathematical geosciences*, 42(1):49–64, 2010. 4
- [14] Richard Hartley and Andrew Zisserman. *Multiple view geometry in computer vision*. Cambridge university press, 2003. 2, 4
- [15] Li Jiang, Hengshuang Zhao, Shaoshuai Shi, Shu Liu, Chi-Wing Fu, and Jiaya Jia. Pointgroup: Dual-set point grouping for 3d instance segmentation. In *Proceedings of the IEEE/CVF conference on computer vision and Pattern recognition*, pages 4867–4876, 2020. 2
- [16] Mrinal Kalakrishnan, Sachin Chitta, Evangelos Theodorou, Peter Pastor, and Stefan Schaal. Stomp: Stochastic trajectory optimization for motion planning. In *2011 IEEE international conference on robotics and automation*, pages 4569–4574. IEEE, 2011. 2, 3
- [17] Jeon Ho Kang, Sagar Joshi, Ruopeng Huang, and Satyandra K Gupta. Robotic compliant object prying using diffusion policy guided by vision and force observations. *IEEE Robotics and Automation Letters*, 2025. 3
- [18] Sertac Karaman and Emilio Frazzoli. Sampling-based algorithms for optimal motion planning. *The international journal of robotics research*, 30(7):846–894, 2011. 3
- [19] Pushkal Katara, Zhou Xian, and Katerina Fragkiadaki. Gen2sim: Scaling up robot learning in simulation with generative models. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 6672–6679. IEEE, 2024. 2
- [20] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Trans. Graph.*, 42(4):139–1, 2023. 2, 6
- [21] John McCormac, Ronald Clark, Michael Bloesch, Andrew Davison, and Stefan Leutenegger. Fusion++: Volumetric object-level slam. In *2018 international conference on 3D vision (3DV)*, pages 32–41. IEEE, 2018. 2
- [22] Ben Mildenhall, Pratul P Srinivasan, Matthew Tancik, Jonathan T Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. *Communications of the ACM*, 65(1):99–106, 2021. 2
- [23] Alexander Millane, Helen Oleynikova, Emilie Wirbel, Remo Steiner, Vikram Ramasamy, David Tingdahl, and Roland Siegwart. nvblox: Gpu-accelerated incremental signed distance field mapping. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 2698–2705. IEEE, 2024. 2, 3, 6
- [24] Mayank Mittal, Calvin Yu, Qinxi Yu, Jingzhou Liu, Nikita Rudin, David Hoeller, Jia Lin Yuan, Ritvik Singh, Yunrong Guo, Hammad Mazhar, et al. Orbit: A unified simulation framework for interactive robot learning environments. *IEEE Robotics and Automation Letters*, 8(6):3740–3747, 2023. 2, 4
- [25] Yao Mu, Tianxing Chen, Zanxin Chen, Shijia Peng, Zhiqian Lan, Zeyu Gao, Zhixuan Liang, Qiaojun Yu, Yude Zou, Mingkun Xu, et al. Robotwin: Dual-arm robot benchmark with generative digital twins. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 27649–27660, 2025. 2
- [26] Adithyavairavan Murali, Balakumar Sundaralingam, Yu-Wei Chao, Wentao Yuan, Jun Yamada, Mark Carlson, 9

- 650 Fabio Ramos, Stan Birchfield, Dieter Fox, and Clemens
651 Eppner. GraspGen: A diffusion-based framework for 6-
652 dof grasping with on-generator training. *arXiv preprint*
653 *arXiv:2507.13097*, 2025. 6
- 654 [27] Kensuke Nakamura, Lasse Peters, and Andrea Bajcsy. Generalizing safety beyond collision-avoidance via latent-space
655 reachability analysis. *arXiv preprint arXiv:2502.00935*,
656 2025. 3
- 657 [28] Richard A Newcombe, Shahram Izadi, Otmar Hilliges,
658 David Molyneaux, David Kim, Andrew J Davison, Pushmeet
659 Kohi, Jamie Shotton, Steve Hodges, and Andrew Fitzgibbon.
660 Kinectfusion: Real-time dense surface mapping and track-
661 ing. In *2011 10th IEEE international symposium on mixed*
662 *and augmented reality*, pages 127–136. Ieee, 2011. 2
- 663 [29] Chuanruo Ning, Kuan Fang, and Wei-Chiu Ma. Prompting
664 with the future: Open-world model predictive control with
665 interactive digital twins. *arXiv preprint arXiv:2506.13761*,
666 2025. 2
- 667 [30] John Oliensis. A critique of structure-from-motion algo-
668 rithms. *Computer Vision and Image Understanding*, 80(2):
669 172–214, 2000. 2
- 670 [31] Onur Özyeşil, Vladislav Voroninski, Ronen Basri, and Amit
671 Singer. A survey of structure from motion*. *Acta Numerica*,
672 26:305–364, 2017. 2
- 673 [32] Jaesik Park, Qian-Yi Zhou, and Vladlen Koltun. Colored
674 point cloud registration revisited. In *Proceedings of the IEEE*
675 *international conference on computer vision*, pages 143–
676 152, 2017. 5
- 677 [33] Nathan Ratliff, Matt Zucker, J Andrew Bagnell, and Sid-
678 dhartha Srinivasa. Chomp: Gradient optimization tech-
679 niques for efficient motion planning. In *2009 IEEE inter-
680 national conference on robotics and automation*, pages 489–
681 494. IEEE, 2009. 2, 3
- 682 [34] Martin Runz, Maud Buffier, and Lourdes Agapito. Maskfu-
683 sion: Real-time recognition, tracking and reconstruction of
684 multiple moving objects. In *2018 IEEE international sym-
685 posium on mixed and augmented reality (ISMAR)*, pages 10–
686 20. IEEE, 2018. 2
- 687 [35] Johannes L Schonberger and Jan-Michael Frahm. Structure-
688 from-motion revisited. In *Proceedings of the IEEE con-
689 ference on computer vision and pattern recognition*, pages
690 4104–4113, 2016. 2, 6
- 691 [36] John Schulman, Jonathan Ho, Alex X Lee, Ibrahim Awwal,
692 Henry Bradlow, and Pieter Abbeel. Finding locally optimal,
693 collision-free trajectories with sequential convex optimiza-
694 tion. In *Robotics: science and systems*, pages 1–10. Berlin,
695 Germany, 2013. 2, 3
- 696 [37] Steven M Seitz, Brian Curless, James Diebel, Daniel
697 Scharstein, and Richard Szeliski. A comparison and evalua-
698 tion of multi-view stereo reconstruction algorithms. In *2006*
699 *IEEE computer society conference on computer vision and*
700 *pattern recognition (CVPR'06)*, pages 519–528. IEEE, 2006.
701 2
- 702 [38] Frank Steinbrücker, Jürgen Sturm, and Daniel Cremers.
703 Real-time visual odometry from dense rgb-d images. In *2011*
704 *IEEE international conference on computer vision work-
705 shops (ICCV Workshops)*, pages 719–722. IEEE, 2011. 2
- 706 [39] Ioan A Sucan, Mark Moll, and Lydia E Kavraki. The open
707 motion planning library. *IEEE Robotics & Automation Mag-
708 azine*, 19(4):72–82, 2012. 3
- 709 [40] Yuzhu Sun, Mien Van, Stephen McIlvanna, Nguyen Minh
710 Nhat, Kabirat Olayemi, Jack Close, and Séán McLoone.
711 Digital twin-driven reinforcement learning for obstacle
712 avoidance in robot manipulators: A self-improving online
713 training framework. *arXiv preprint arXiv:2403.13090*, 2024.
714 2
- 715 [41] Balakumar Sundaralingam, Siva Kumar Sastry Hari, Adam
716 Fishman, Caelan Garrett, Karl Van Wyk, Valts Blukis,
717 Alexander Millane, Helen Oleynikova, Ankur Handa, Fabio
718 Ramos, et al. Curobo: Parallelized collision-free robot mo-
719 tion generation. In *2023 IEEE International Conference on*
720 *Robotics and Automation (ICRA)*, pages 8112–8119. IEEE,
721 2023. 2, 4
- 722 [42] Kenta Tanaka. cupoch – robotics with gpu computing, 2020.
723 <https://github.com/neka-nat/cupoch>. 2, 5
- 724 [43] Matthew Tancik, Ethan Weber, Evonne Ng, Rui long Li,
725 Brent Yi, Terrance Wang, Alexander Kristoffersen, Jake
726 Austin, Kamyar Salahi, Abhik Ahuja, et al. Nerfstudio: A
727 modular framework for neural radiance field development.
728 In *ACM SIGGRAPH 2023 conference proceedings*, pages 1–
729 12, 2023. 6
- 730 [44] Gabriel Taubin. A signal processing approach to fair surface
731 design. In *Proceedings of the 22nd annual conference on*
732 *Computer graphics and interactive techniques*, pages 351–
733 358, 1995. 5
- 734 [45] Marcel Torne, Anthony Simeonov, Zechu Li, April Chan,
735 Tao Chen, Abhishek Gupta, and Pukit Agrawal. Reconciling
736 reality through simulation: A real-to-sim-to-real approach
737 for robust manipulation. *arXiv preprint arXiv:2403.03949*,
738 2024. 2
- 739 [46] Jianyuan Wang, Minghao Chen, Nikita Karaev, Andrea
740 Vedaldi, Christian Rupprecht, and David Novotny. Vggt: Vi-
741 sual geometry grounded transformer. In *Proceedings of the*
742 *Computer Vision and Pattern Recognition Conference*, pages
743 5294–5306, 2025. 2, 4
- 744 [47] Weiyue Wang, Ronald Yu, Qiangui Huang, and Ulrich Neu-
745 mann. Sgn: Similarity group proposal network for 3d
746 point cloud instance segmentation. In *Proceedings of the*
747 *IEEE conference on computer vision and pattern recogni-
748 tion*, pages 2569–2578, 2018. 2
- 749 [48] Lai Wei, Jiahua Ma, Yibo Hu, and Ruimao Zhang. Ensuring
750 force safety in vision-guided robotic manipulation via im-
751 plicit tactile calibration. *arXiv preprint arXiv:2412.10349*,
752 2024. 3
- 753 [49] Xinyue Wei, Minghua Liu, Zhan Ling, and Hao Su. Approx-
754 imate convex decomposition for 3d meshes with collision-
755 aware concavity and tree search. *ACM Transactions on*
756 *Graphics (TOG)*, 41(4):1–18, 2022. 5
- 757 [50] Yuxuan Wu, Lei Pan, Wenhua Wu, Guangming Wang, Yanzi
758 Miao, Fan Xu, and Hesheng Wang. Rl-gsbridge: 3d gaus-
759 sian splatting based real2sim2real method for robotic ma-
760 nipulation learning. In *2025 IEEE International Conference*
761 *on Robotics and Automation (ICRA)*, pages 192–198. IEEE,
762 2025. 2

- 764 [51] Jianyun Xu, Song Wang, Ziqian Ni, Chunyong Hu, Sheng
765 Yang, Jianke Zhu, and Qiang Li. Sam4d: Segment
766 anything in camera and lidar streams. *arXiv preprint*
767 *arXiv:2506.21547*, 2025. 2
- 768 [52] Mutian Xu, Xingyilang Yin, Lingteng Qiu, Yang Liu, Xin
769 Tong, and Xiaoguang Han. Samprom3d: Locating sam
770 prompts in 3d for zero-shot instance segmentation. In *2025*
771 *International Conference on 3D Vision (3DV)*, pages 1222–
772 1232. IEEE, 2025. 2
- 773 [53] Jiahui Yang, Jason Jingzhou Liu, Yulong Li, Youssef Khaky,
774 Kenneth Shaw, and Deepak Pathak. Deep reactive policy:
775 Learning reactive manipulator motion planning for dynamic
776 environments. *arXiv preprint arXiv:2509.06953*, 2025. 2, 3
- 777 [54] Yao Yao, Zixin Luo, Shiwei Li, Tian Fang, and Long Quan.
778 Mvsnet: Depth inference for unstructured multi-view stereo.
779 In *Proceedings of the European conference on computer vi-*
780 *sion (ECCV)*, pages 767–783, 2018. 2
- 781 [55] Haoqi Yuan, Ziye Huang, Ye Wang, Chuan Mao, Chaoyi
782 Xu, and Zongqing Lu. Demograsp: Universal dexterous
783 grasping from a single demonstration. *arXiv preprint*
784 *arXiv:2509.22149*, 2025. 2
- 785 [56] Dingyuan Zhang, Dingkang Liang, Hongcheng Yang,
786 Zhikang Zou, Xiaoqing Ye, Zhe Liu, and Xiang Bai. Sam3d:
787 Zero-shot 3d object detection via segment anything model.
788 *arXiv preprint arXiv:2306.02245*, 2023. 2
- 789 [57] Yuchen Zhou, Jiayuan Gu, Tung Yen Chiang, Fanbo Xiang,
790 and Hao Su. Point-sam: Promptable 3d segmentation model
791 for point clouds. *arXiv preprint arXiv:2406.17741*, 2024. 2
- 792 [58] Zhengxue Zhou, Xingyu Yang, Hao Wang, and
793 Xuping Zhang. Digital twin with integrated robot-
794 human/environment interaction dynamics for an industrial
795 mobile manipulator. In *2022 International Conference on*
796 *Robotics and Automation (ICRA)*, pages 5041–5047. IEEE,
797 2022. 2